

CSC 369

Introduction to Distributed Computing

Distributing Data

Two (complementary) approaches to distributing data across nodes:

- Replication
 - Keep an exact copy of data on multiple nodes
- Partitioning
 - Split a large dataset into smaller subsets, store each subset on a different node

Separate mechanisms, but often used together

Partitioning

When data or request volume becomes too large for a single node, we must break up the data into **partitions**



In this course, we will use the term "partition"
Other terminology you may encounter:

- shard (MongoDB, Elasticsearch)
- region (HBase)
- tablet (BigTable)
- vnode (Cassandra, Riak)
- vBucket (Couchbase)

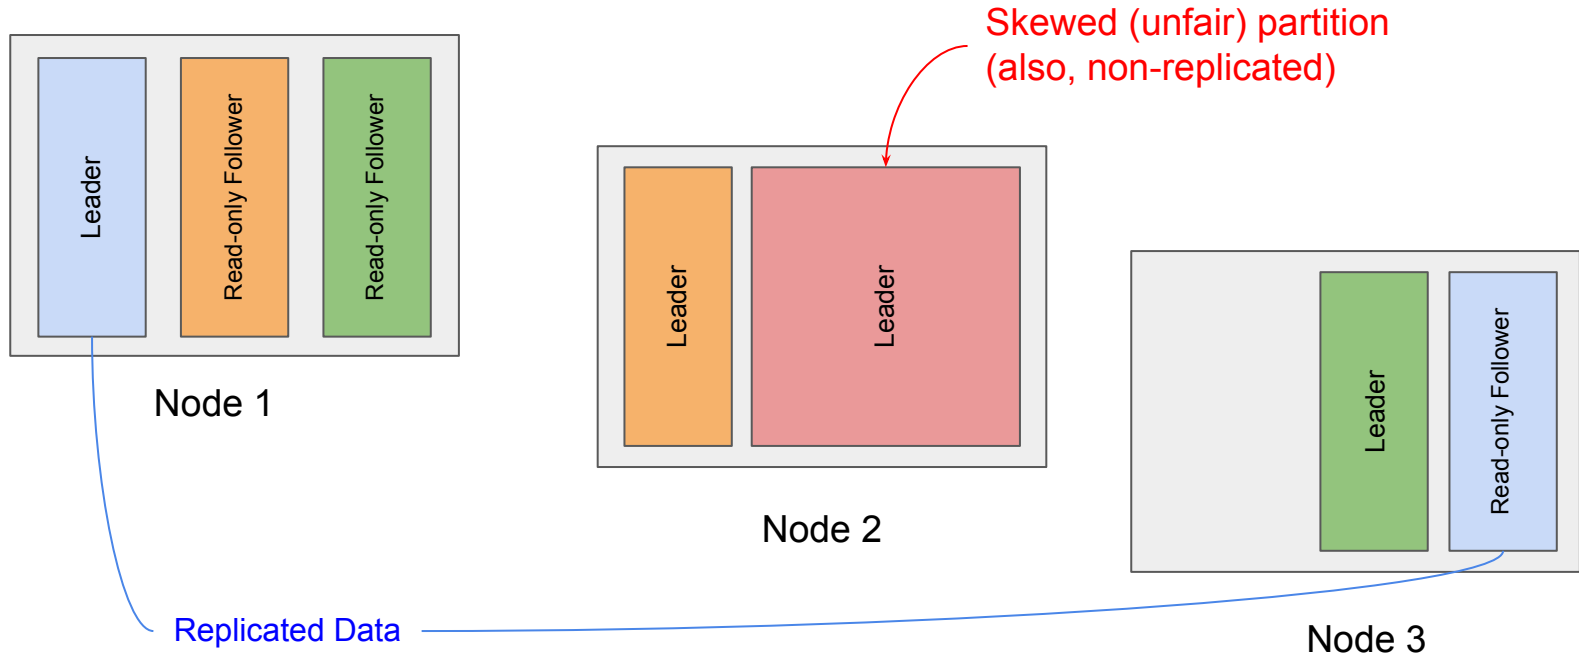
Partitioning

- Each piece of data (record, row, document) belongs to exactly *one* partition
- Different partitions placed on different nodes
- For *operations on single partitions*, each node independently operates on its own

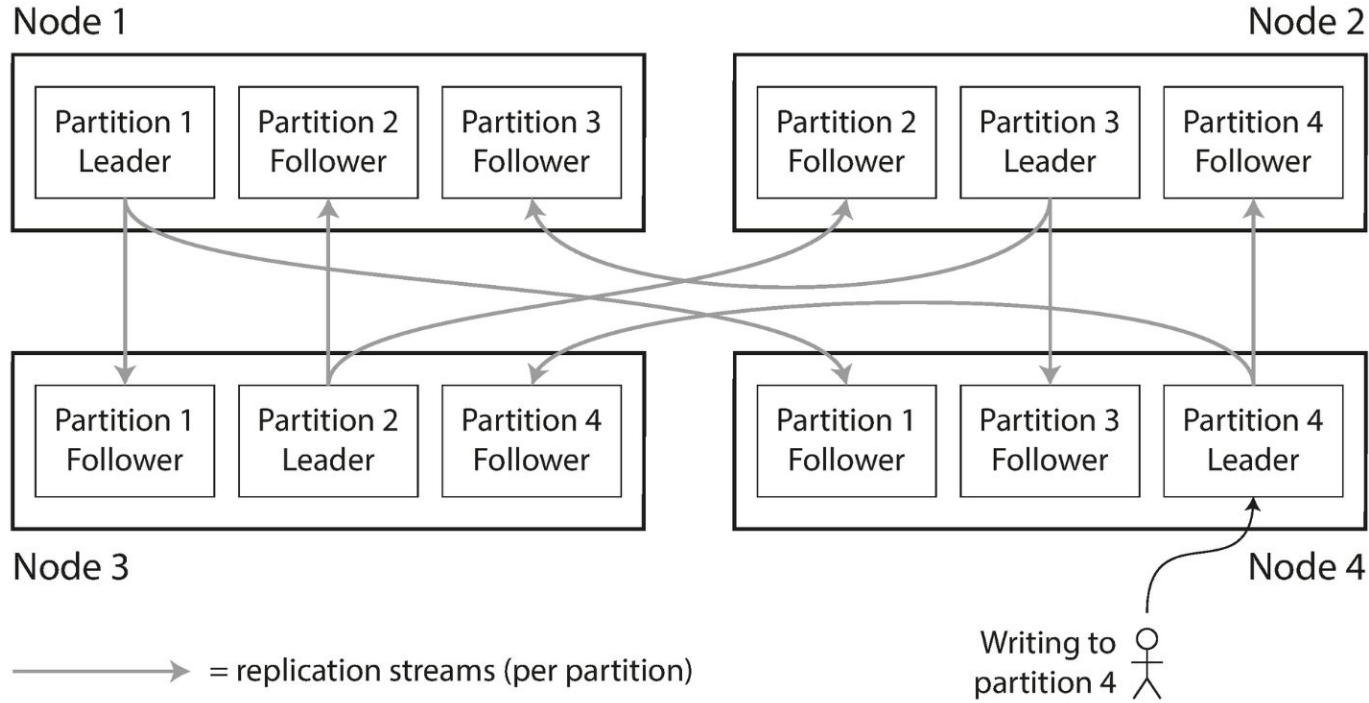
Partitioning Considerations

- Complex operations that must span multiple partitions
- Addition/removal of nodes
- Rebalancing partitions
- Routing requests to the correct partitions
- Even distribution of data and load

Partitioning and Replication



Partitioning and Replication



Partitioning Data

Goal: spread data and query load evenly

Strategies:

- Randomly assign data to partitions
- Assign a range to each partition
- Hash-based
- Parent-child interleaving


Partitioning Strategy - Random

Assign records to partitions records completely at random

Pros	Cons
simple	look up because it's random
even if do it correctly	

Partitioning Strategy - Range-Based

Assume each record has a key. Assign a continuous range of keys to each partition. ex. year, name, etc.



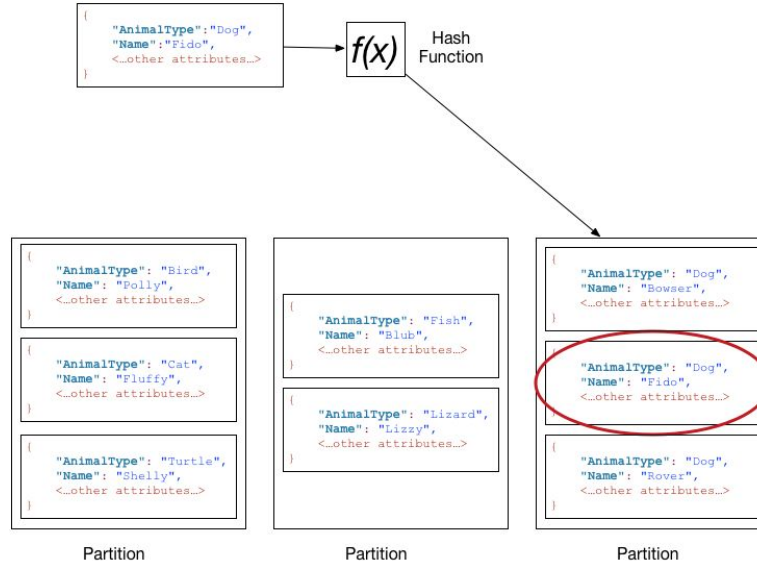
Partitioning Strategy - Range-Based

Assume each record has a key. Assign a continuous range of keys to each partition.

Pros	Cons
common data will be next to each other	might overload node if don't know the range a head of time. example, some keys are more popular than others.
instantaneous look up by key name	

Partitioning Strategy - Hash-Based

Assume each record has a key. Apply a hash function to each key, assign each partition a range of hashes.




Partitioning Strategy - Hash-Based

Assume each record has a key. Apply a hash function to each key, assign each partition a range of hashes.

Pros	Cons
look up is constant	if the hash function is bad, it might lead to over loaded nodes
	search is not efficient
	in a range of data

Partitioning by Hash of Key



Partitioning Strategy - Interleaved

????????

Co-locate parent/child records (eg. user & user's events), to minimize cost of commonly-performed joins.

```
{ id: 123, email: "user@host" }  
{ user: 123, event: "comment", tsamp: "..." }  
{ user: 123, event: "delete", tsamp: "..." }  
{ user: 123, event: "create", tsamp: "..." }  
{ id: 234, email: "user234@host" }  
{ user: 234, event: "comment", tsamp: "..." }  
...
```

Partition 1

```
{ id: 456, email: "user456@host" }  
{ user: 456, event: "create", tsamp: "..." }  
{ user: 456, event: "comment", tsamp: "..." }  
{ user: 456, event: "delete", tsamp: "..." }  
{ id: 567, email: "user567@host" }  
{ user: 567, event: "create", tsamp: "..." }  
{ user: 567, event: "create", tsamp: "..." }  
...
```

Partition 2

Partitioning and Indexes

- (Simplistic) assumption thus far: all lookups are by key
- Support for secondary indexes often useful
 - For example: find cars by color/make/model
 - Key feature of relational databases, search tools (Elasticsearch, Solr)
- Problem: how to map secondary indexes to partitions?

Secondary Index - Partitioned by Document

independence. write is very efficient

Partition 0

PRIMARY KEY INDEX	
191	→ {color:"red", make:"Honda", location:"Palo Alto"}
214	→ {color:"black", make:"Dodge", location:"San Jose"}
306	→ {color:"red", make:"Ford", location:"Sunnyvale"}

SECONDARY INDEXES (Partitioned by document)	
color:black	→ [214]
color:red	→ [191, 306]
color:yellow	→ []
make:Dodge	→ [214]
make:Ford	→ [306]
make:Honda	→ [191]

Partition 1

PRIMARY KEY INDEX	
515	→ {color:"silver", make:"Ford", location:"Milpitas"}
768	→ {color:"red", make:"Volvo", location:"Cupertino"}
893	→ {color:"silver", make:"Audi", location:"Santa Clara"}

SECONDARY INDEXES (Partitioned by document)	
color:black	→ []
color:red	→ [768]
color:silver	→ [515, 893]
make:Audi	→ [893]
make:Ford	→ [515]
make:Volvo	→ [768]

scatter/gather read from all partitions



"I am looking for a red car"

Secondary Index, Partitioned by Term

a write is written to multiple partition. it is not independence anymore

Partition 0

PRIMARY KEY INDEX

191 → {color: "red", make: "Honda", location: "Palo Alto"}
214 → {color: "black", make: "Dodge", location: "San Jose"}
306 → {color: "red", make: "Ford", location: "Sunnyvale"}

SECONDARY INDEXES (Partitioned by term)

color:black → [214]
color:red → [191, 306, 768]
make:Audi → [893]
make:Dodge → [214]
make:Ford → [306, 515]

Partition 1

PRIMARY KEY INDEX

515 → {color: "silver", make: "Ford", location: "Milpitas"}
768 → {color: "red", make: "Volvo", location: "Cupertino"}
893 → {color: "silver", make: "Audi", location: "Santa Clara"}

SECONDARY INDEXES (Partitioned by term)

color:silver → [515, 893]
color:yellow → []
make:Honda → [191]
make:Volvo → [768]

Text



"I am looking for a red car"

Rebalancing Partitions

- Query load grows, need more CPUs
- Data volume grows, need more disk/RAM
- Node fails
- Need to move data from one node to another, a process called **rebalancing**

Rebalancing Considerations

- After rebalancing, load & data should be evenly distributed among nodes
- Uninterrupted availability during rebalancing process
- Minimize amount of data moved between nodes

Rebalancing Strategies

- $\text{hash}(\text{key}) \bmod N$ (where N is the number of nodes)
- For example: 10 nodes, mod 10 returns 0-9
 - Assign to appropriate node
- Good strategy? assume even size and a good hash function


data transfer is not efficient because we have to move the whole thing

Rebalancing: Fixed Number of Partitions

- Create large number of partitions (many more than node count)
- Assign several partitions to each node
- New nodes can "steal" partitions from existing nodes
- If nodes are removed, partitions are divided between remaining nodes

Rebalancing Example

Before rebalancing (4 nodes in cluster)



After rebalancing (5 nodes in cluster)

Legend:

----- partition remains on the same node

————> partition migrated to another node

How Many Partitions?

- If each partitions holds a large amount of data, rebalancing and recovery are expensive
- If partitions hold very little data, overhead is high maintenance become a nightmare
- Ideal: "just right"

based on the situation to decide where the balance is

Dynamic Partitioning

- Fixed number of partitions / fixed boundaries not always convenient
- Some implementations begin with a single partition
- As a partition grows, it is split
 - One half transferred to another node
- Partition count adapts to data volume

good and common strategy

Rebalancing: Automatic or Manual?

- Decision to split or move partitions
 - Fully automatic? fully manual? somewhere in between?
- Rebalancing can be expensive
 - Need to take care to avoid cascading failures

cascading effect if something gone wrong even when the system is taking care fo the rebalancing
- Human-assisted process can avoid issues
 - up the user to commit

Request Routing

How do we route client requests to the appropriate node?

1. Clients connect to any node, request forwarded as needed
2. Routing tier (partition-aware load balancer)
3. Clients aware of partitioning configuration

?????????

Key issue: propagating change information (new partitions, splits)

All participants must agree

Replication / Partitioning Summary

- Replication and Partitioning
 - Distribute data across nodes for scaling, latency, or availability reasons
 - Two separate approaches, often used together
- Automatic or Manual?
- Managing changing data (writes)
- Increase/decrease number of nodes