

CSC 369

Introduction to Distributed Computing

Filtering

Projection
Transformation

Grouping

Aggregation

Join

Sort

Ungrouping
Unwinding

~~Limit~~

~~Skip~~

Sample

Join in MapReduce

R



S



Join condition

$$R.x == S.y$$

Join in MapReduce

R

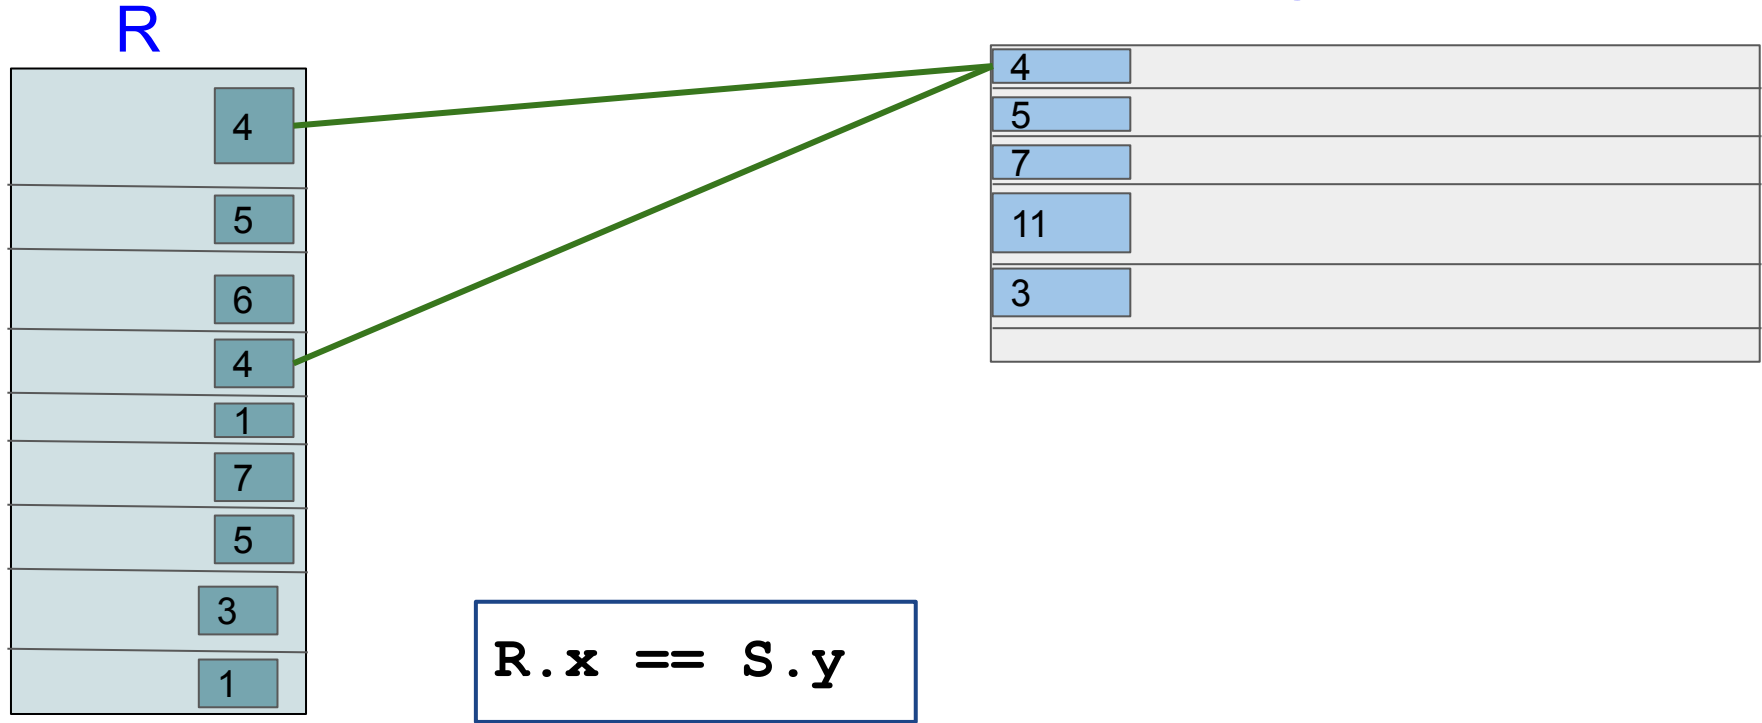
	4
	5
	6
	4
	1
	7
	5
	3
	1

S

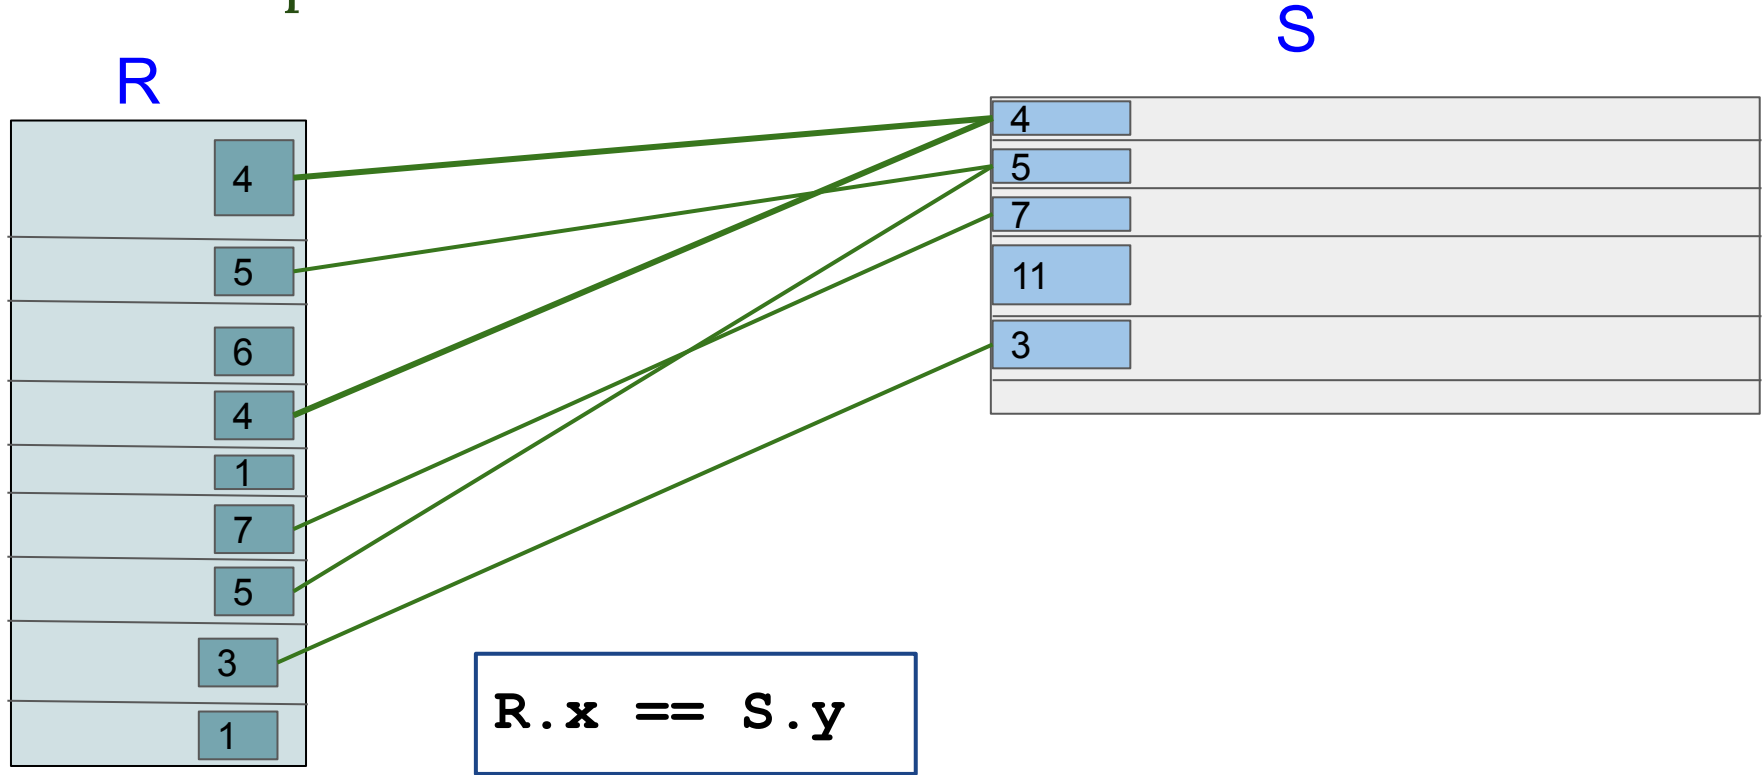
4	
5	
7	
11	
3	

`R.x == S.y`

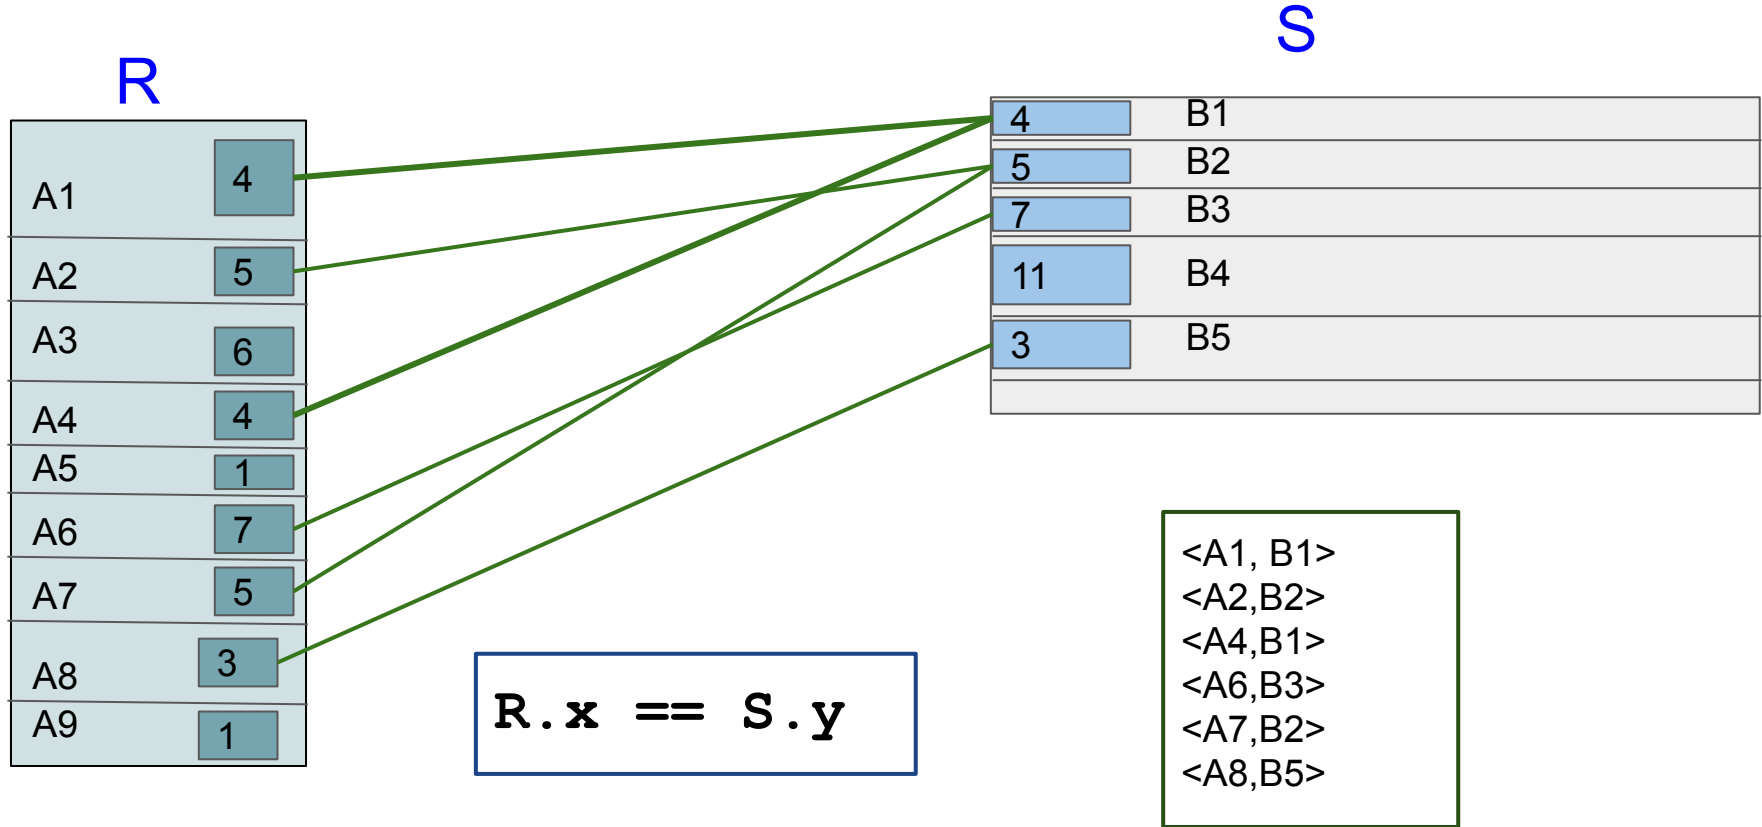
Join in MapReduce



Join in MapReduce



Join in MapReduce



Join in MapReduce

Straightforward implementation: nested loop

```
for each  $t$  in  $R$  do
  for each  $s$  in  $S$  do
    If  $t.x == s.y$  then emit( $\langle t, s \rangle$ )
```


Join in MapReduce

Straightforward implementation: nested loop

```
for each t in R do
  for each s in S do
    If t.x == s.y then emit(<t,s>)
```


Need to turn this into distributed computation

Join in MapReduce

Technical Issue To Solve: **work with multiple files**


Join in MapReduce

Technical Issue To Solve: **work with multiple files**



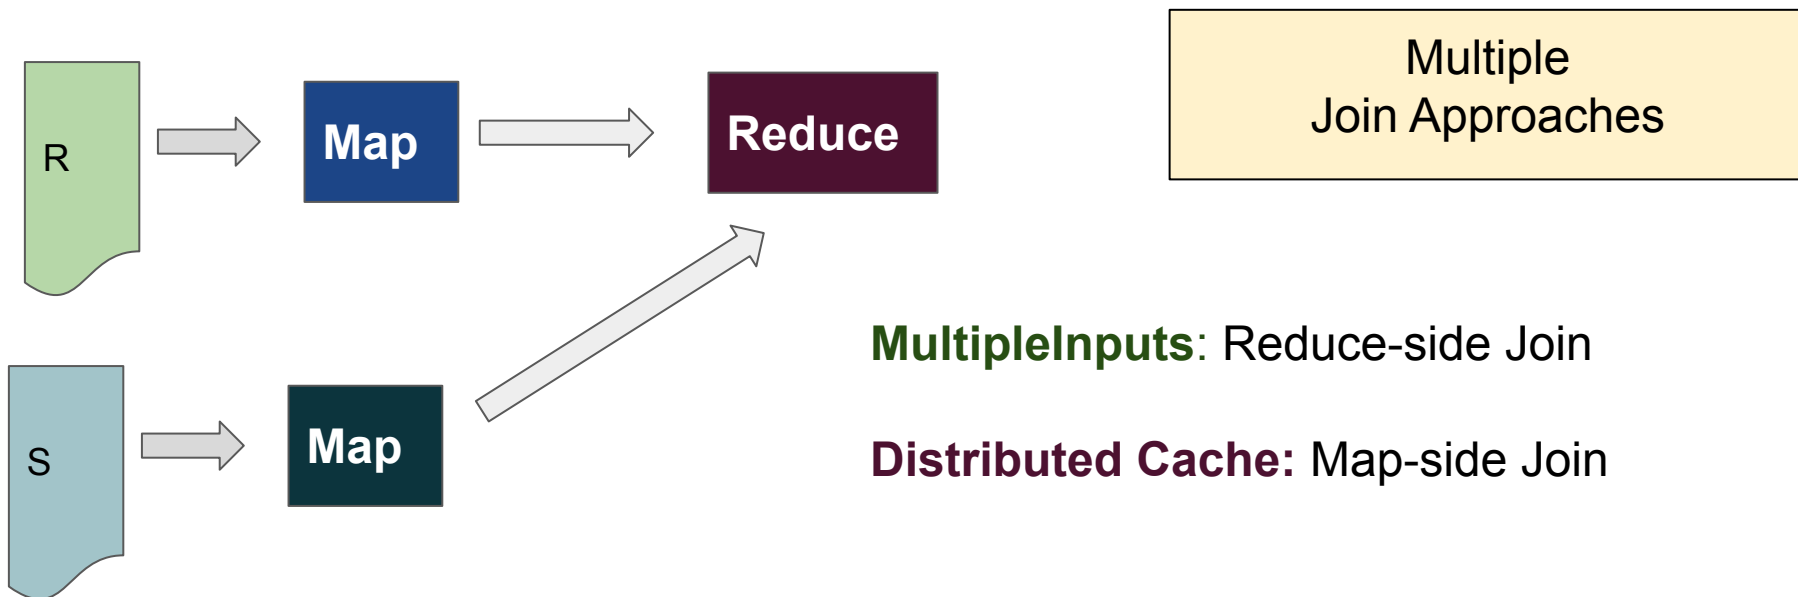
Join in MapReduce

Technical Issue To Solve: **work with multiple files**



Join in MapReduce

Technical Issue To Solve: **work with multiple files**




Reduce-side Join

Idea: process two sources independently, “in parallel”

Merge records in **reduce**


Reduce-side Join

$$R.x == S.y$$



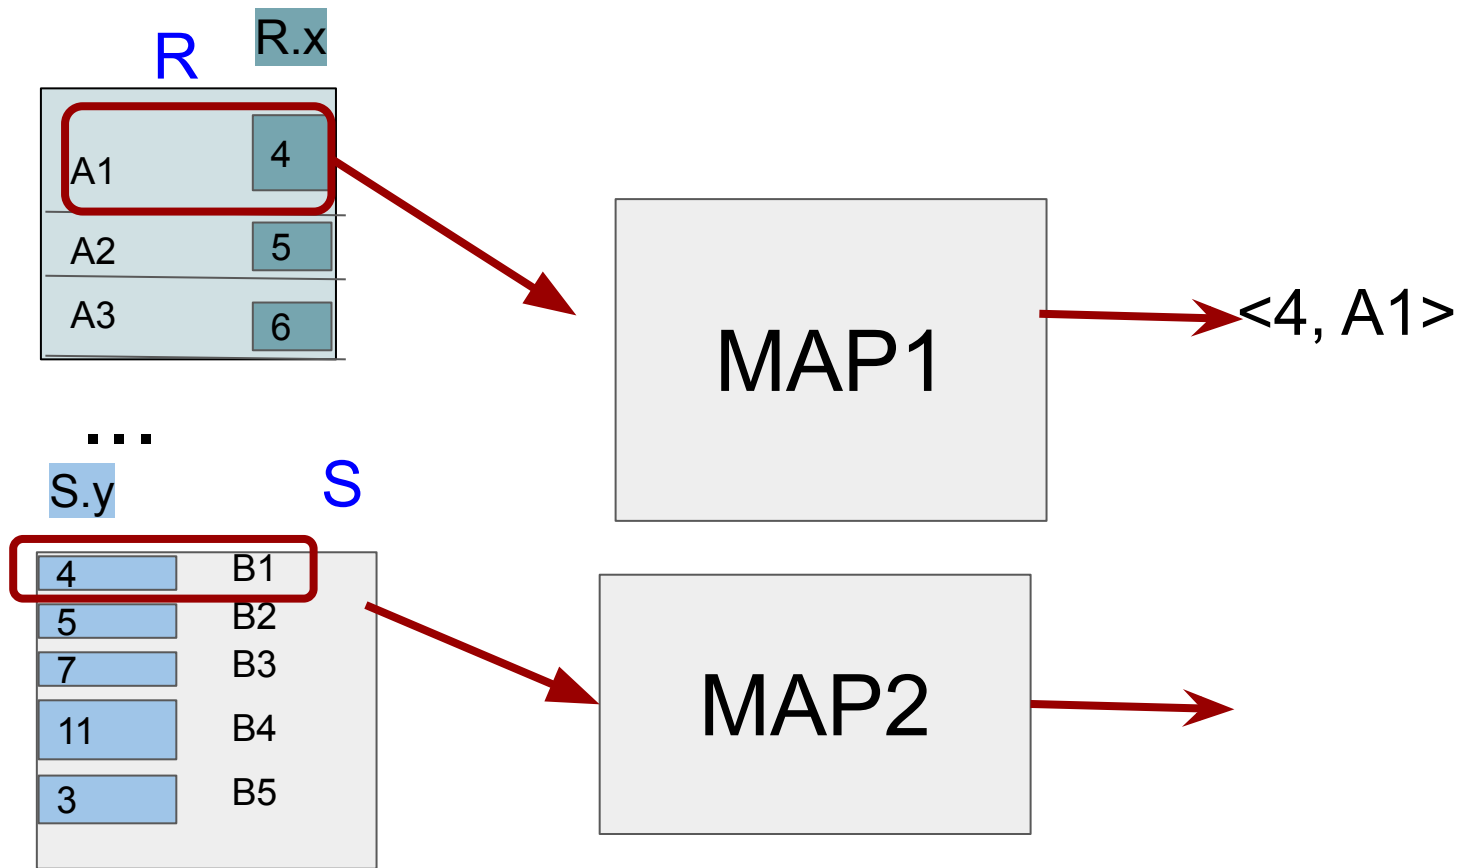
Reduce-side Join

$R.x == S.y$




Reduce-side Join

$R.x == S.y$




Reduce-side Join

$R.x == S.y$




Reduce-side Join

R.x == S.y




Reduce-side Join

$$R.x == S.y$$



Reduce-side Join

$$R.x == S.y$$



Reduce-side Join


$\langle 4, (A1, A4, B1) \rangle$

Need a Trick!!!

Include Source Information into the output value of Map

Reduce-side Join

$$R.x == S.y$$



Reduce-Side Join: Maps

```
mapR(key, value)
  newKey = value.x
  newValue={"source":1,
           "Content":value}
  emit(newKey,newValue)
```


Reduce-Side Join: Maps

Encode the source

```
mapR(key, value)
  newKey = value.x
  newValue={"source":1,
            "Content":value}
  emit(newKey,newValue)
```

```
mapS(key, value)
  newKey = value.y
  newValue={"source":2,
            "Content":value}
  emit(newKey,newValue)
```

Reduce-Side Join: Maps + Reduce

```
mapR(key, value)
  newKey = value.x
  newValue={"source":1,
            "Content":value}
  emit(newKey,newValue)
```

```
mapS(key, value)
  newKey = value.y
  newValue={"source":2,
            "Content":value}
  emit(newKey,newValue)
```

```
reduce(key, Iterable values)
  fromR = []
  fromS = []
  for v in values do
    if v.source == 1 then
      fromR.append(v)
    else
      fromS.append(v)
  end for
```

Step 1: break input list apart by source

Reduce-Side Join: Maps + Reduce

```
mapR(key, value)
  newKey = value.x
  newValue={"source":1,
            "Content":value}
  emit(newKey,newValue)
```

```
mapS(key, value)
  newKey = value.y
  newValue={"source":2,
            "Content":value}
  emit(newKey,newValue)
```

```
reduce(key, Iterable values)
  fromR = []
  fromS = []
  for v in values do
    if v.source == 1 then
      fromR.append(v)
    else
      fromS.append(v)
    end for
```

Step 1: break input list apart by source

Reduce-Side Join: Maps + Reduce

```
mapR(key, value)
  newKey = value.x
  newValue={"source":1,
            "Content":value}
  emit(newKey,newValue)
```

```
mapS(key, value)
  newKey = value.y
  newValue={"source":2,
            "Content":value}
  emit(newKey,newValue)
```

```
reduce(key, Iterable values)
```

```
  fromR = []
  fromS = []
  for v in values do
    if v.source == 1 then
      fromR.append(v)
    else
      fromS.append(v)
    end for
  for r in fromR do
    for s in fromS do
      emit(key, (r,s))
    end for
  end for
```

Step 2: emit cartesian product of records from R cross records from S

Reduce Side Join Demo (UserMessages)

Reduce-Side Joins

Advantages	Disadvantages
<p>Combine two sources -- any size</p> <p>Sorting by join attribute</p>	<p>$(R.x == S.y) = \text{Empty Set}$</p> <p>If one user is extremely chatty -- most of the work could be performed by 1 reducer</p> <p>Sorting by join attribute</p> <p>Imbalance in size of inputs</p>

Map-Side Join

R

A1	4
A2	5
A3	6
A4	4
A5	1
A6	7
A7	5
A8	3
A9	1

Assumptions:

R is LARGE

S is SMALL (fits main memory)

Map-Side Join

R

A1	4
A2	5
A3	6
A4	4
A5	1
A6	7
A7	5
A8	3
A9	1

Assumptions:

R is LARGE

S is SMALL (fits in main memory)

Idea:

Split R

Load S onto every compute node

Map-Side Join

R

A1	4
A2	5
A3	6
A4	4
A5	1
A6	7
A7	5
A8	3
A9	1

S

4	B1
5	B2
7	B3
11	B4
3	B5

S

4	B1
5	B2
7	B3
11	B4
3	B5

Map-Side Join

R

A1	4
A2	5
A3	6
A4	4
A5	1
A6	7
A7	5
A8	3
A9	1

Assumptions:

R is LARGE

S is SMALL (fits main memory)

Idea:

Split R

Load S onto every compute node

Map-Side Join

Distributed Cache

R

A1	4
A2	5
A3	6
A4	4
A5	1
A6	7
A7	5
A8	3
A9	1

S

4	B1
5	B2
7	B3
11	B4
3	B5

S

4	B1
5	B2
7	B3
11	B4
3	B5

Map-side Join Mapper:

```
HashMap dCache      // distributed cache is an instance
                    // variable in Mapper class

setup(cacheFile):
    // upload cache file into a convenient data structure
    file = open(cacheFile, "read")

    for each line in cacheFile do
        record = parse(line)      // use whatever parsing needed
        dCache.insert(record)
    end for
```

Map-Side Join Mapper:

```
HashMap dCache      // distributed cache is an instance
                    // variable in Mapper class

map(key, value):    // map gets input parameters from the
                    // other file

// naive loop implementation, improve with hash or sort
for record in dCache do
    if join condition on the pair (value, record) is true
    then emit(null, (value, record))
```

Map-Side Join Mapper:

```
reduce(key, Iterable values):
```

```
    for record in values do
```

```
        emit(null, values)
```

```
    end for
```

Map-Side Joins

Advantages	Disadvantages
<p data-bbox="112 366 620 405">Just one mapper to implement</p> <p data-bbox="112 456 475 494">No extra sorting effort</p>	<p data-bbox="1000 366 1653 405">Not suitable when both inputs are large</p> <p data-bbox="1000 456 1224 494">No "free" sort</p>